

Meridian

AI agents run HR and finance — with human oversight

Pre-seed · Peter Alesso, Founder · September 2026

enterprise-hr-finance.vercel.app

ACT ONE

The problem

Enterprises have been handed a new AI worker, but how do we hold it accountable?

Almost everyone has bought AI. Almost nobody has shipped.

79%

have adopted AI agents

11%

run them in production

Gartner expects 40% of agentic projects to be cancelled by 2027. The reason it gives is **weak governance**.

Sources: enterprise AI agent adoption surveys, 2026; Gartner, 2026.

The money leaves, and you find out sixty days later.

A controller approves **6,000 invoices a month**. She already knows some are duplicates.

She learns which ones from an outside audit firm that keeps a share of what it recovers.

OFFERED TODAY

- An assistant she must supervise
- A chat box with no record of who decided
- A pilot her auditors reject

She does not need a smarter model. She needs a name to hold accountable.

Duplicate payments run 0.1% to 2% of AP spend. Sources: APQC Open Standards Benchmarking; AP recovery audit benchmarks, 2026.

ACT TWO

How we solve it

Narrow agents that do the work, and a governance layer that makes them accountable.

Agents do the work. A question box stops at a human.

Twelve agents work

Each scoped to one workflow, running continuously.



A person decides

Anything consequential waits, with the decision recorded against a name.

Enforced below the model, not written into a prompt. **That distinction is the business.**

Twelve agents. Each owns one job and one number.

Help Desk HR cases resolved	Recruiting Candidates screened	Payroll Errors caught pre-run	Scheduling Open shifts filled
Performance Reviews drafted	Job Architecture Pay bands benchmarked	Audit Evidence packaged	Planning Variances explained
Controls Transactions tested	Close Month-end orchestrated	Revenue Contracts Revenue risk flagged	Contract Review Paper redlined

Nobody says yes to a general-purpose agent with access to payroll. **They say yes to one workflow.**

We stop the loss before it happens, not after.

	RECOVERY AUDIT TODAY	MERIDIAN
When	After the money has left	Before the payment run
Cycle	45 to 90 days	Continuous
Coverage	A sample	Every transaction

\$100M of AP spend carries \$100K to \$500K of annual leakage. Firms are already paid a share to find it late. One agent pays for the platform.

Priced in credits per completed action, not per seat.

Arithmetic on published duplicate-payment rates, not a customer result.

ACT THREE

How we should be funded

What is built in the next eighteen months.

Built and deployed. No revenue, no customers, no team.

DONE

- The commercial platform, live and public
- Twelve agents specified end to end
- Real-time operations layer, in production

NOT DONE — THIS IS THE RAISE

- The execution layer. The dashboard is a simulation
- No database, no SSO, no stored Registry
- Zero revenue, zero customers, zero pilots

Design partners shown in the product are illustrative, not real companies. **We will not present them as traction.**

THE ASK

\$1.5M pre-seed. Eighteen months of runway.

\$1.5M

pre-seed round

18

months of runway

3

engineering hires

3

paying customers, verified

MONTHS 0–6

3 paid pilots live

MONTHS 6–12

Verified saving at 3 customers

MONTHS 12–18

Series A metrics, or stop

62% to engineering. The rest to pilot delivery, SOC 2 Type II and infrastructure. What we want besides money: one introduction to a CFO who will run a paid pilot.

Agents will do the work. Somebody still has to be accountable.

Meridian is the record of who that is — for every agent an enterprise runs, including the ones we did not build.

enterprise-hr-finance.vercel.app

Peter Alesso · Founder